Module 01: Elliptic Curve Cryptography

Week-03: Cryptanalysis of ECDSA

Jan Gilcher, Kenny Paterson, and Fernando Virdia

jan.gilcher@inf.ethz.ch, kenny.paterson@inf.ethz.ch,
fernando.virdia@inf.ethz.ch

September 2021

Hi all! Welcome to the first Information Security Lab for this semester. This lab is part of the Module-01 on *Elliptic Curve Cryptography*, or ECC in short. Due to technical limitations the labs will only take place in presence. You are also welcome to post your questions and to join ongoing discussions on the Moodle forum at the following link:

`https://moodle-app2.let.ethz.ch/course/view.php?id=15448`

Students who cannot attend the labs in presence are especially encouraged to use Moodle to ask questions.

1 Overview

This lab is focused on implementing cryptanalysis of ECDSA [1] based on *nonce-leakages*. In other words, you will be implementing attack algorithms that recover the secret ECDSA signing key based on either complete or partial information about the nonce(s) used in the ECDSA signing algorithm. We will look at three flavors of cryptanalytic attacks, namely *known-nonce attacks*, *repeated-nonce attacks* and *partially-known-nonce attacks*. The first two attacks are relatively straightforward and serve as warm-ups, while the third one is significantly more involved and will require the use of lattice-based techniques. You will implement two variants of the third attack, one where the most significant bits of the nonce are leaked, the other where the least significant bits of the nonce are leaked. Finally, you will also implement the *partial-known-nonce* attacks on Schnorr signatures [5], a similar signature scheme to ECDSA that will be described in this lab sheet.

You have already seen the theoretical descriptions of these attacks during the lectures and in the exercise classes; but we will recap the necessary material in this lab sheet to help with the implementation. Please note that we expect you to use Python-3 for this lab. We will not accept submissions/solutions coded in any other language, including older versions of Python. *So please make sure that your submissions are Python-3 compatible.*

2   *fpylll and other required libraries*

This lab requires the use of solvers for lattice problems such as the closest vector problem (CVP) and the shortest vector problem (SVP). Fortunately, we can make black-box use of publicly available implementations of these (and many other) lattice algorithms from `fpylll` - a Python library for performing lattice reduction on lattices over the integers, available at:

`https://github.com/fplll/fpylll`

`fpylll` is a Python wrapper around its the C++ library `fplll` [1] and provides an easy-to-use interface for several state-of-the-art algorithms on lattices that rely on floating-point computations. This includes implementations of lattice reduction algorithms such as LLL [2] and BKZ [3], as well as implementations of solvers for CVP and SVP, all of which will be relevant and extremely useful for this lab. Hence, we expect you to familiarize yourself with `fpylll`.

[1] `https://github.com/fplll/fplll`

Note that installing `fpylll` can be a time-consuming process. It requires prior installation of several other open-source libraries for arbitrary precision integer and floating-point arithmetic, as well as libraries for interfacing between Python and C++. While you are welcome to try and install `fpylll` on your own work-stations, we are also providing a virtual machine (VM) image that has `fpylll` pre-installed and ready to use.

`fpylll` is not the only required library for this lab. In Sections 6+ you may also want to use the `pyca/cryptography` library. A third optional (but enabled by default) library is `tqdm`, which is used in the test files and allows you to keep track of the estimated completion time for each exercise when running your solution. Both `pyca/cryptography` and `tqdm` are provided inside the VM.

Note to Apple M1 users: We have not been able to test all possible virtualization solutions running on Apple silicon. If your preferred solution fails, you may want to consider using CoCalc instead of running the VM (see next Section).

The link to access and download the VM is available on Moodle. Below we provide you with a simple set of instructions on how to use this VM, and how to run `fpylll` on it:

- Boot the VM using your favorite virtualization software. For example, you can download and use the Oracle VM VirtualBox from the url below:[2]

  `https://www.virtualbox.org/`

- Login using the user `isl` and password `isl`. If needed, the root password is also set to `isl`.

- `fpylll`, `pyca/cryptography` and `tqdm` are installed system wide. Launching a terminal and running the `python3` or `ipython3` interpreters should allow you to successfully `import fpylll, cryptography, tqdm`.

[2] VMWare Workstation Player should also successfully load the VM, after relaxing virtual hardware compliance checks when importing it (you should be prompted).

- If importing the libraries works, you can now run any Python3 script(s) for solving the lab.

Note that when automatically evaluating your code, we will use this same VM and the same build environment. So using the VM would additionally allow you to pre-test your submission in an environment resembling our automated testing environment.

Using CoCalc

In case you are facing issues with installing the VM and running any of the required libraries on it, you can alternatively test your code online using CoCalc. Please see the relevant instructions below:

- Go to `https://cocalc.com`.

- Click on "Run CoCalc Now" (or sign in with an existing account).

- **(Optional)** Create an account if you want to save your work on the CoCalc servers.

- Once the page has loaded, click on "Files" (upper left).

- On the upper right click "Upload" and select the files containing the input test vectors that are provided to you, and the test file `module_1_ECDSA_Cryptanalysis_tests.py`.

There are at least two ways you can now proceed.

Option 1:

- On the upper left click on "+ New" and then click on "Jupyter Notebook". Under "Suggested Kernels", choose "SageMath". The Jupyter notebook should be created and will automatically open.

- Copy and paste your code into the notebook.

- Select the cell containing your code and hit `shift+enter`.

- Your code will now run (as indicated by filled a green circle). This might take up to 10 minutes (maybe more depending on the amount of load on the Cocalc servers).

- Please note that CoCalc may interrupt and stop your code if you are idle for too long. So make sure to keep an eye on it!

Option 2:

- Upload also your script containing the solution to the lab.

- In your list of files, double click on the solution script to open it.

- Here, you can make any modifications you desire.

- To run the solution file, open the "Code" drop down menu on the top-right of the screen, then choose "Terminal".

- A shell will open on the same directory as your files. To run your solution file (say, "sol.py"), use the SageMath interpreter by calling `sage sol.py`.

Note: The reason we suggest using SageMath instead of Python3 when running on CoCalc is that SageMath is essentially a Python3 distribution including all the libraries required for this lab (and much more![3]).

3   Getting Started

You are provided with a skeleton file named `module_1_ECDSA_Cryptanalysis_Skel.py`. This file has several unimplemented functions related to cryptanalysis of ECDSA that you are expected to implement. Since this is the second lab, you are expected to implement any low-level functions that you wish to use on your own. The skeleton file defines a few such functions but does not implement them:

- `egcd`: Computes the gcd of two integers using the Euclidean algorithm.

- `mod_inv`: Computes $a^{-1} \mod p$ for input integers a and p.

4   ECDSA Cryptanalysis: Single Known Nonce

We now describe the first warm-up task for this lab – implementing ECDSA cryptanalysis based on a single known nonce. The skeleton file has an (incomplete) implementation of `recover_x_known_nonce` – a function that is meant to realize this attack. You are expected implement this function. It takes as input:

- k: the nonce used by the ECDSA signing algorithm (represented as a big integer)

- h: the message on which the signature was produced, post-hashing and mapping to an integer modulo q (represented as a big integer)

- (r, s): the signature produced by the ECDSA signing algorithm (represented as a pair of big integers)

- q: the order of the base point P on the elliptic curve used by the ECDSA scheme (represented as a big integer)

and should output the recovered signing key x (represented as a big integer).

Note: This attack is quite straightforward to implement and involves simple arithmetic modulo q. For the relevant theory for this attack, look at the lecture slides.

5 ECDSA Cryptanalysis: Repeated Nonces

The second warm-up task for this lab involves implementing ECDSA cryptanalysis based on a pair of signatures on *distinct* messages sharing a *common* (but *unknown*) nonce. In particular, the skeleton file has an (incomplete) implementation of `recover_x_repeated_nonce` – a function that is meant to realize this attack. You are expected implement this function. It takes as input:

- (h_1, h_2): the messages on which the signatures were produced, post-hashing and mapping to integers modulo q (represented as a pair of big integers)

- (r_1, s_1): the signature produced by the ECDSA signing algorithm on the first message (represented as a pair of big integers)

- (r_2, s_2): the signature produced by the ECDSA signing algorithm on the second message (represented as a pair of big integers)

- q: the order of the base point P on the elliptic curve used by the ECDSA scheme (represented as a big integer)

and should output the recovered signing key x (represented as a big integer).

Note: Again, this attack is quite straightforward to implement and involves simple linear algebra modulo q. The relevant theory for this attack can be found in the lecture slides.

6 ECDSA Cryptanalysis: Partially Known Nonces

We now arrive at the main task for this lab – implementing ECDSA cryptanalysis based on partially known nonces using lattice reduction algorithms. For the ease of implementation, we have broken up the attack into a sequence of modular tasks. For each task, the skeleton file describes an incomplete function that you are expected to complete.

We will now walk you through setting up the attack against ECDSA when the most significant bits (MSBs) of the nonce are leaked. You will then have to repeat the analysis yourself for dealing with the case where the least significant bits (LSBs) of the nonce are leaked instead.

Note that some of the functions described below will take as input the flags "`algorithm`" and "`givenbits`". During testing and marking of this exercise,

these will be set to "ecdsa" and "msbs" respectively, to indicate to your functions that one is running the attack against ECDSA when the MSBs of the nonce are leaked. In later exercises (see Sections 7 and 8), you will implement the other code branches corresponding to the different possible values of these flags.

A Single HNP Equation

The first sub-task is to implement a function that builds a single equation as part of a larger overall *hidden number problem* (HNP) instance. The equation is to be built from a single ECDSA signature and some partial information about the corresponding nonce used by the ECDSA signing algorithm.

More concretely, assume that you are provided with the L most significant bits of an N-bit nonce used to generate a signature (r, s) on a message h (post-hashing and mapping to an integer modulo q – this is the representation of a message we will use in the rest of the discussion). Recall from the lecture notes that:

$$s = k^{-1} \cdot (h + x \cdot r) \mod q,$$

which after some simple re-arrangement yields:

$$(r \cdot s^{-1}) \cdot x = k - h \cdot s^{-1} \mod q.$$

Setting $t = (r \cdot s^{-1}) \mod q$ and $z = (h \cdot s^{-1}) \mod q$, we have:

$$tx = k - z \mod q.$$

Now, let a be an integer represented by the L most significant bits of k. Then we can write:

$$k = a \cdot 2^{N-L} + 2^{N-L-1} + e,$$

where e is some *error* term bounded as:

$$0 \leq |e| \leq 2^{N-L-1}.$$

Setting $u = a \cdot 2^{N-L} + 2^{N-L-1} - z$, we have:

$$tx = u + e \mod q$$

Here t and u are known, x is the target unknown and e is small but unknown. We can think of computing (t, u) as setting up a single equation as part of a larger overall HNP instance.

The skeleton file has an (incomplete) implementation of `setup_hnp_single_sample` – a function that is meant to realize this sub-task. You are expected to implement this function. It takes as input:

- N: the total number of bits in the nonce k;

- L: the number of known most significant bits of the nonce k;

- $(k_1, \ldots, k_L)$: a Python list object containing the L most (or least, depending on the exercise) significant bits of the nonce k (k_1 being the most significant bit of the leak);

- h: the message on which the signature was produced, post-hashing and mapping to an integer modulo q;

- (r, s): the signature produced by the signing algorithm;

- q: the order of the base point P on the elliptic curve used by the signature scheme;

- givenbits, a flag set to "msbs" or "lsbs" depending on whether we are given the most or the least significant bits of k;

- algorithm, a flag set to "ecdsa" or "ecschnorr" depending on whether we are constructing a HNP instance from an ECDSA signature (this exercise) or an EC-Schnorr signature (see last exercise);

and should output the pair (t, u) (to be represented as a pair of big integers).

For ease and modularity of implementation, you can implement and use the subroutine MSB_to_Padded_Int, which takes as input N, L and $(k_1, \ldots, k_L)$ as described above, and outputs $a \cdot 2^{N-L} + 2^{N-L-1}$.

We have provided you with some test vectors to test that your implementation of the function setup_hnp_single_sample is consistent with our expectation (e.g., with respect to the manner in which the bits of k are ordered in the input). The parameters for the test vectors are $(N, L) = (256, 128)$. They are in the file testvectors_hnp_single_sample_256_128.txt.

While we do not separately evaluate your implementation of this function, its correctness is crucial for the overall correctness of your attack implementation; so we strongly encourage you to test this function against the test vectors provided to you.

Building the Overall HNP Instance

The second sub-task is to simply extend the aforementioned approach to build a larger HNP instance consisting of several equations, built from several ECDSA signatures and partial leakages from the corresponding nonces. The skeleton file has an (incomplete) implementation of setup_hnp_all_samples – a function that is meant to realize this sub-task. You are expected implement this function. It takes as input:

- N: the total number of bits in each nonce k_i;

- L: the number of known most significant bits of each nonce k_i;

- n: the number of equations to be built (called n in the lecture);

- $\{(k_{i,1}, \ldots, k_{i,L})\}_{i \in [n]}$: a Python nested list object containing L most significant bits of each nonce k_i;

- $\{h_i\}_{i \in [n]}$: a Python list object containing each message h_i post-hashing and mapping to an integer modulo q;

- $\{(r_i, s_i)\}_{i \in [n]}$: a pair of Python list objects containing each signature produced by the signing algorithm;

- q: the order of the base point P on the elliptic curve used by the ECDSA scheme;

- `givenbits`, a flag set to `"msbs"` or `"lsbs"` depending on whether we are given the most or the least significant bits of k;

- `algorithm`, a flag set to `"ecdsa"` or `"ecschnorr"` depending on whether we are constructing a HNP instance from an ECDSA signature (this exercise) or an EC-Schnorr signature (see last exercise);

and should output $\{(t_i, u_i)\}_{i \in [n]}$ (to be represented as a pair of Python list objects). Obviously, you should iterate over each signature and use the `setup_hnp_single_sample` function you implemented earlier to compute the corresponding (t_i, u_i) pair.

We have also provided you with some test vectors to test that your implementation of the function `setup_hnp_all_samples` is consistent with our expectation. The parameters for the test vectors are $(N, L, n) = (256, 128, 5)$. They are in the file `testvectors_hnp_all_samples_256_128_5.txt`.

Again, while we do not separately evaluate your implementation of this function, its correctness is crucial for the overall correctness of your attack implementation; so we strongly encourage you to test this function against the test vectors provided to you.

HNP to CVP

The next step is to transform the HNP instance into an instance of the closest vector problem (CVP). Recall from the lecture notes that given $\left(\{(t_i, u_i)\}_{i \in [n]}, q \right)$, one can construct the following CVP basis matrix:

$$B_{\text{CVP}} = \begin{bmatrix} q & 0 & 0 & \ldots & 0 & 0 \\ 0 & q & 0 & \ldots & 0 & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \ldots & q & 0 \\ t_1 & t_2 & t_3 & \ldots & t_n & 1/2^{L+1} \end{bmatrix},$$

and the following CVP target vector:

$$\underline{u}_{\text{CVP}} = \begin{bmatrix} u_1 & u_2 & u_3 & \ldots & u_n & 0 \end{bmatrix}.$$

As was formally established during the lecture, for appropriate choices of parameters N, L, n and q, given B_{CVP} and $\underline{u}_{\text{CVP}}$, a CVP solver should produce a

vector

$$\underline{v} = \begin{bmatrix} v_1 & v_2 & v_3 & \ldots & v_n & v_{n+1} \end{bmatrix}.$$

such that v_{n+1} is of the form:

$$v_{(n+1)} = x/2^{L+1}.$$

Question 1: Suppose that the CVP output vector contains $(x - q)/2^{L+1}$ instead of $x/2^{L+1}$. Could this occur?

Hint: Recall that our CVP instance $(B_{\text{CVP}}, \underline{u}_{\text{CVP}})$ looks like the following:

$$B_{\text{CVP}} = \begin{bmatrix} q & 0 & 0 & \ldots & 0 & 0 \\ 0 & q & 0 & \ldots & 0 & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \ldots & q & 0 \\ t_1 & t_2 & t_3 & \ldots & t_n & 1/2^{L+1} \end{bmatrix} , \quad \underline{u}_{\text{CVP}} = \begin{bmatrix} u_1 & u_2 & u_3 & \ldots & u_n & 0 \end{bmatrix},$$

As already discussed, we expect that given this instance, the CVP solver should output a vector $\underline{v}$ in the lattice generated by B_{CVP} that is "close" to $\underline{u}_{\text{CVP}}$ and is of the form:

$$\underline{v} = \begin{bmatrix} v_1 & v_2 & v_3 & \ldots & v_n & x/2^{L+1} \end{bmatrix}.$$

Note that we must have:

$$\underline{v} = \begin{bmatrix} a_1 & a_2 & a_3 & \ldots & a_n & x \end{bmatrix} B_{\text{CVP}}.$$

for some $a_1, \ldots, a_n \in \mathbb{Z}$ (in fact $a_i = -\ell_i$ where the ℓ_i satisfy $t_i x = u_i + e_i + \ell_i q$ as in the lecture). Now, consider the following modified vector

$$\underline{v}' = \begin{bmatrix} (t_1 + a_1) & (t_2 + a_2) & (t_3 + a_3) & \ldots & (t_n + a_n) & (x - q) \end{bmatrix} B_{\text{CVP}}$$

Quite evidently, $\underline{v}'$ is also a vector in the lattice generated by B_{CVP}. Additionally, convince yourself that the following is true:

$$\underline{v}' = \begin{bmatrix} v_1 & v_2 & v_3 & \ldots & v_n & (x - q)/2^{L+1} \end{bmatrix}.$$

Given this, do you think that the following condition could hold?

$$||\underline{v}' - \underline{u}_{\text{CVP}}|| < ||\underline{v} - \underline{u}_{\text{CVP}}||.$$

If yes, the CVP algorithm would output $\underline{v}'$ instead of $\underline{v}$, and your implementation would recover $(x - q)$ instead of q.

Question 2: The basis matrix described above has a non-integral entry. Figure out if the `fpylll` in-built CVP solver supports basis matrices and/or target vectors with non-integral entries.

Question 3: Suppose that the `fpylll` in-built CVP solver *does not* support basis matrices and/or target vectors with non-integral entries. How would you transform the aforementioned CVP instance into one where the basis matrix has only integral entries?

Hint: Think about scaling both the basis matrix and the target vector with an appropriate scalar.

To summarize, in this sub-task, you are expected to implement the incomplete function `hnp_to_cvp` that takes as input:

- N: the total number of bits in each nonce k_i;

- L: the number of known most significant bits of each nonce k_i;

- n: the number of equations in the HNP instance;

- $\{(t_i, u_i)\}_{i \in [n]}$: a pair of Python list objects representing the HNP instance;

- q: the order of the base point P on the elliptic curve used by the ECDSA scheme;

and that outputs $(B_{\text{CVP}}, \underline{u}_{\text{CVP}})$, where B_{CVP} is the CVP basis matrix (to be represented as a Python nested list object) and $\underline{u}_{\text{CVP}}$ is the CVP target vector (to be represented as a Python list object). Both of these should be scaled appropriately to contain only integral entries (if required).

Solving CVP using `fpylll`

The next natural sub-task is to implement a function that solves the (potentially scaled) CVP instance using the `fpylll` in-built CVP solver. So go ahead and implement the function `solve_cvp` in the skeleton file. This function takes as input $(B_{\text{CVP}}, \underline{u}_{\text{CVP}})$ in the format as described above and outputs the solution vector $\underline{v}$ (represented as a Python list object).

Question 4: Do you want to perform some pre-processing on the basis matrix B_{CVP} before passing it as input to the `fpylll` in-built CVP solver? Does the `fpylll` documentation say something about this?

Hint: You might want to pre-process the basis matrix B_{CVP} by reducing it using some lattice reduction algorithms. In particular, the `fpylll` library has in-built functions implementing lattice reduction algorithms such as LLL and BKZ.

Question 5: If you do not immediately find a satisfactory answer to Question 4, here is a more directed question. Would pre-processing B_{CVP} using a lattice reduction algorithm help? If yes, why? Think about what in-built reduction algorithms in the `fpylll` library you could use here.

CVP to SVP via Kannan's Embedding

Alternatively, suppose that you do not have direct access to a CVP solver, but you have access to a solver for the shortest vector problem (SVP). As shown in the lecture, one can further convert the CVP instance into an SVP instance, and then solve the resulting SVP instance. In particular, given the (potentially scaled) CVP basis B_{CVP} and the CVP target vector $\underline{u}_{\text{CVP}}$, once can use Kannan's embedding [4] to prepare the following lattice basis matrix:

$$B_{\text{SVP}} = \begin{bmatrix} B_{\text{CVP}} & \mathbf{0} \\ \underline{u}_{\text{CVP}} & M \end{bmatrix},$$

where $\mathbf{0}$ is an all-zero column vector of appropriate dimension and M is a specially chosen constant.

As was formally established during the lecture, for an appropriate choice of M, the lattice generated by B_{SVP} as designed above should contain a "short" vector of the form

$$\underline{f}' = \begin{bmatrix} \underline{f} & M \end{bmatrix},$$

where $\underline{f} = \underline{u}_{\text{CVP}} - \underline{v}$, and $\underline{u}_{\text{CVP}}$ is the target CVP vector, while $\underline{v}$ is the target CVP solution. Note that given $\underline{f}'$, one can extract immediately recover $\underline{f}$, and hence the CVP solution $\underline{v}$.

Question 6: From what you have seen in the lecture, what should your choice of M be assuming that B_{CVP} and $\underline{u}_{\text{CVP}}$ are *not* scaled?

Hint: Think about the maximum value that M could take, and how that relates to the maximum distance between the CVP target and solution vectors.

Question 7: Now suppose that B_{CVP} and and $\underline{u}_{\text{CVP}}$ *are* scaled. How should you scale M to preserve SVP correctness?

To summarize, in this sub-task, you are expected to implement the incomplete function `cvp_to_svp` that takes as input a CVP instance $(B_{\text{CVP}}, \underline{u}_{\text{CVP}})$ and outputs the Kannan-embedded basis B_{SVP} (to be represented as a Python nested list object).

Note: If the CVP instance $(B_{\text{CVP}}, \underline{u}_{\text{CVP}})$ is scaled appropriately to contain only integral entries, then B_{SVP} also contains only integral entries.

Solving SVP using `fpylll`

Again, the next natural sub-task is to implement a function that solves the SVP instance using the `fpylll` in-built SVP solver. So go ahead and implement the function `solve_svp` in the skeleton file. This function takes as input B_{SVP} in the format as described above and should output a lattice basis containing the solution vector $\underline{f}'$ (represented as a Python list object) as described above. Note that

in the next step you will implement `recover_x_partial_nonce_SVP`. This function will need to identify the location in the basis of the desired solution vector $\underline{f}'$, and extract the signing key x from it.

Practical cryptanalysis often needs to deal with heuristics and unlikely corner cases, for example regarding the expected location in the basis of $\underline{f}'$. In order to get full marks in this and following exercises, you will have to carefully handle possible corner cases. To this aim it may be useful to develop your own test vectors for the current exercise and the following ones.

Question 8: Would pre-processing B_{SVP} using a lattice reduction algorithm help? Or is this redundant when solving SVP? What does the `fpylll` in-built SVP solver actually do?

Question 9: For the specially structured basis matrix B_{SVP} prepared using Kannan's embedding, is the desired solution vector $\underline{f}' = \begin{bmatrix} \underline{f} & M \end{bmatrix}$ always the shortest vector in the lattice generated by B_{SVP}? Could the lattice contain a vector shorter than $\underline{f}'$? If yes, what could such a vector be?

Hint: Revisit the exercises for this week, recalling that the lattice generated by B_{SVP} contains the following vector:

$$\begin{bmatrix} 0 & 0 & 0 & \dots & q/2^{L+1} & 0 \end{bmatrix},$$

or an appropriately scaled version of the same. Is this vector shorter than the desired solution vector $\underline{f}' = \begin{bmatrix} \underline{f} & M \end{bmatrix}$?

Question 10: Suppose that after some experimentation, you managed to convince yourself that in practice, with overwhelmingly large probability, the desired solution vector $\underline{f}'$ is not the shortest vector but the *second* shortest vector in the lattice generated by B_{SVP}. In this case, does the `fpylll` in-built SVP solver allow you to recover the second shortest vector instead of the shortest vector?

Remark: In light of Questions 9 and 10, it could be appropriate to check if the value x' recovered during the attack inside function `recover_x_partial_nonce_SVP` (see the next step) is indeed the same as the signing key x. This can be done by defining a function `check_x`. This function takes as input

- x': the candidate value of x recovered using the SVP solver;

- (Q_x, Q_y): a tuple containing the x and y coordinates of the verification key $Q = [x]P$;

and returns `True` if $x' = x$ and `False` otherwise. We provide this function in the skeleton file for this lab, implemented using the `pyca/cryptography` Python library.[4]

Putting Everything Together

Finally, you are expected to put everything together by implementing the partially completed functions `recover_x_partial_nonce_CVP` and `recover_x_partial_nonce_SVP`. These functions are expected to implement partial nonce-based ECDSA cryptanalysis by directly solving CVP and solving CVP via SVP, respectively. Both these functions take as input the following:

- (Q_x, Q_y): a tuple containing the x and y coordinates of the verification key $Q = [x]P$;

- N: the total number of bits in each nonce k_i;

- L: the number of known most significant bits of each nonce k_i;

- n: the number of equations to be built;

- $\{(k_{i,1}, \ldots, k_{i,L})\}_{i \in [n]}$: a Python nested list object containing L most significant bits of each nonce k_i;

- $\{h_i\}_{i \in [n]}$: a Python list object containing each message h_i post-hashing and mapping to an integer modulo q;

- $\{(r_i, s_i)\}_{i \in [n]}$: a pair of Python list objects containing each signature produced by the signing algorithm;

- q: the order of the base point P on the elliptic curve used by the signature scheme;

- `givenbits`, a flag set to `"msbs"` or `"lsbs"` depending on whether we are given the most or the least significant bits of k;

- `algorithm`, a flag set to `"ecdsa"` or `"ecschnorr"` depending on whether we are constructing a HNP instance from an ECDSA signature (this exercise) or an EC-Schnorr signature (see last exercise);

and should output the recovered signing key x (represented as a big integer). The functions should use the sub-routines that you have already implemented before. Some of this is already done for you in the skeleton file. The missing piece is how to finally extract the signing key x, which you should implement, thereby completing the functions.

Recall from the lecture notes that solving the CVP instance should output a vector $\underline{v}$ of length $(n + 1)$ such that the last entry of $\underline{v}$ is $x/2^{L+1}$.

Question 11: Recall that you might have already scaled your CVP basis and target vector in order to use the `fpylll` built-in CVP solvers. In this case, what would you expect the solution to the CVP instance to look like?

7 ECDSA Cryptanalysis: Leaking Least Significant Bits

Having implemented the attack against ECDSA when the most significant bits of the nonce are leaked in Section 6, you will now augment your code to handle the case where the least significant bits of the nonce are leaked instead. This case was not explicitly covered in class, however it can also be reduced to the Hidden Number Problem, and solved using the same lattice reduction techniques.

In particular, you will be required to modify `setup_hnp_single_sample`, `setup_hnp_all_samples`, `recover_x_partial_nonce_CVP` and `recover_x_partial_nonce_SVP` to handle the case where the input flag `algorithm` is set to "ecdsa" and the flag `givenbits` is set to "lsbs".

The core modification required will be the construction of t and u inside `setup_hnp_single_sample`. In Section 6, this was done starting with the equation used to derive the term s in the signature,

$$s = k^{-1} \cdot (h + x \cdot r) \mod q, \tag{1}$$

and observing the particular form of k when the L most significant bits were known. Now, let $\hat{a}$ be an integer representing the leaked L least significant bits of k. To setup the HNP in this case, observe that k is of the form

$$k = e \cdot 2^L + \hat{a},$$

where e is some *error* term bounded as:

$$0 \le |e| < q/2^L.$$

Think about how to rearrange Equation (1) to turn it into an equation of the form

$$tx = u + e \mod q$$

where t and u are known, which can be solved using the functions implemented in the previous exercise.

For ease of implementation, you can use the subroutine `LSB_to_Int`, which takes as input $(k_1, \ldots, k_L)$ (in this exercise, the L *least* significant bits of k), and outputs $\hat{a}$.

Note that in Section 6 you were given test vectors for `setup_hnp_single_sample` and `setup_hnp_all_samples`. You will not be given these now. However, you will be given test vectors for `recover_x_partial_nonce_CVP` and `recover_x_partial_nonce_SVP` (see Section 9).

8 Schnorr Signatures Cryptanalysis: Partially Known Nonces

In Sections 6 and 7, we considered the problem of solving a HNP instance derived from ECDSA signatures where either the least or most significant bits of the nonce have been leaked. However, ECDSA is not the only signature scheme where this

attack strategy is viable. Indeed, the DSA signature scheme from which ECDSA is derived was introduced as a variant of Schnorr signature scheme [5]. An elliptic curve version of Schnorr signature scheme, "EC-Schnorr" can similarly be attacked when enough adjacent bits of a nonce k are leaked.

EC-Schnorr shares the same key generation algorithm as ECDSA: given a public elliptic curve point P, sample a signing key x uniformly at random in $\{1, \ldots, q-1\}$ and publish the verification key $Q = [x]P$.

In order to Sign a message m using the signing key x, one does the following:

1. Select the nonce k uniformly at random from $\{1, \ldots, q-1\}$;

2. Compute $r = \text{x-coord}([k]P) \bmod q$;

3. Compute $h = \text{bits2int}(H(m, r)) \bmod q$, where H is a predetermined cryptographic hash function such as SHA-256;

4. Compute $s = k - h \cdot x \bmod q$.

The signature consists of the pair (h, s).

To verify a signature (h, s) on a message m using the public key Q,

1. Compute $r' = \text{x-coord}([s]P + [h]Q)$;

2. Compute $h' = \text{bits2int}(H(m, r')) \bmod q$;

3. If $h = h'$, then output 1, else output 0.

In this exercise, you will be required to further augment the previously implemented functions in order to recover the signing key x from a list of signatures $\{(h_i, s_i)\}_{i \in [n]}$, where for each signature you will be provided with a list of bits from the nonce that were leaked $\{(k_{i,1}, \ldots, k_{i,L})\}_{i \in [n]}$. As in Sections 6 and 7, you will be required to provide solutions for the case where the L leaked bits are respectively the most or the least significant of the nonce. The elliptic curve used will be the same used for ECDSA, meaning that the modulus q will be the same.

In order to implement the attack, you will be again required to modify `setup_hnp_single_sample`, `setup_hnp_all_samples`, `recover_x_partial_nonce_CVP` and `recover_x_partial_nonce_SVP` to handle the case where the input flag `algorithm` is set to "ecschnorr" and the flag `givenbits` is set to either "msbs" or "lsbs".

The core modification will once again be to `setup_hnp_single_sample`, where you will be required to perform your own analysis to construct the HNP samples (t_i, u_i) by starting from the (h_i, s_i) pairs and the leaked bits $(k_{i,1}, \ldots, k_{i,L})$.

Note that in Section 6 you were given test vectors for `setup_hnp_single_sample` and `setup_hnp_all_samples`. You will not be given these now. However, you will be given test vectors for `recover_x_partial_nonce_CVP` and `recover_x_partial_nonce_SVP` (see Section 9).

9 Testing and Evaluation

Together with the skeleton file, we have provided a test script including four test modules – all of which involve file reads and writes. You can use these modules to test your implementation. The test modules are summarized as follows (in each case the base point order q is a 256-bit prime):

1. The first module tests the correctness of the `recover_x_known_nonce` function over 50 different randomly generated signing keys, nonces and signatures.

2. The second module tests the correctness of `recover_x_repeated_nonce` function over 50 different randomly generated signing keys, nonces and signature pairs sharing the same nonce.

3. The two remaining modules test the correctness of two functions, namely the `recover_x_partial_nonce_CVP` and `recover_x_partial_nonce_SVP` functions over 50 different randomly generated signing keys, nonces and signatures for every possible value of the `algorithm` and `givenbits` flags. For each signing key, the test vectors are generated using the following parameters:

   (a) $(N, L, n) = (256, 128, 5)$
   (b) $(N, L, n) = (256, 32, 10)$
   (c) $(N, L, n) = (256, 16, 20)$
   (d) $(N, L, n) = (256, 8, 60)$

   In other words, the test modules offer increasingly smaller amounts of leakage on the nonce (smaller L) but compensate by providing you with more instances per signing key (larger n).

   In total,

   $$|\{\text{CVP}, \text{SVP}\} \times \{\text{ECDSA}, \text{EC-Schnorr}\} \times \{\text{LSBs}, \text{MSBs}\} \times \{(a), (b), (c), (d)\}| = 32$$

   tests are run on `recover_x_partial_nonce_CVP` and `recover_x_partial_nonce_SVP`, one per every challenge. As mentioned above, each of these combinations will be run on 50 inputs during testing and marking.

For each of these test modules, you are provided with the input test vectors and the corresponding output vectors in separate input and output files. When you execute each test module, you will generate an output file. Your implementation is correct if the contents of this file *exactly* match the contents of the output file provided to you. This check is performed automatically by the tests modules, and a warning is raised if the check fails.

To run the tests, you need to make the following function call to `run_tests`. This code is already included at the end of the skeleton file.

```
from module_1_ECDSA_Cryptanalysis_tests import run_tests
run_tests(recover_x_known_nonce,
    recover_x_repeated_nonce,
    recover_x_partial_nonce_CVP,
    recover_x_partial_nonce_SVP)
```

Note: You are free to test your code using your own custom-designed test modules. However, we will be using the *same* test modules as in the skeleton file to evaluate your submissions under an automated evaluation framework. Your final submission should have these test modules as is. *Please do not tamper with the test modules in any way to avoid interfering with our automated evaluation frameworks.*

Evaluation

When we evaluate your submissions, we will run the exact same test modules, albeit with respect to our own privately generated input and output files, which will not be made public prior to evaluation. This is why we ask you to leave the test module codes as is in your final submissions. *Failure to adhere to this instruction will incur heavy penalties.*

Summary of Evaluation Criteria. To summarize, you will be evaluated based on the correctness of your implementation of the individual functions:

1. `recover_x_known_nonce` (3 points)

2. `recover_x_repeated_nonce` (3 points)

3. `recover_x_partial_nonce_CVP` (3 points for each of the 8 input combinations on ECDSA, 1 point for each of the 8 input combinations on EC-Schnorr)

4. `recover_x_partial_nonce_SVP` (3 points for each of the 8 input combinations on ECDSA, 1 point for each of the 8 input combinations on EC-Schnorr)

So a total of 70 points is available for this week's lab.

For each exercise, points will only be awarded if all the problem instances are successfully solved. For example, if `recover_x_partial_nonce_SVP` is being tested on an ECDSA, MSBs set of instances with $L = 128$, all 50 secret keys used during marking have to be recovered in order to obtain the 3 points. If even one of the 50 secret keys is not recovered, no points will be awarded instead.

Submission Format

Your completed submission for this week should consist of a *single* Python file, and should be named "`module_1_ECDSA_Cryptanalysis.py`".

You are expected to upload your submission to Moodle. The submissions for weeks 2 and 3 should be bundled into a single archive file named "module_1_submission_[insert LegiNo].zip". Submission instructions for the solution to week 2 have already been provided separately in the lab sheet for week 2.

In conclusion, *happy coding*!

10 References

1. *Public Key Cryptography For The Financial Services Industry: Agreement Of Symmetric Keys Using Discrete Logarithm Cryptography.* ANSI X9.42-2003 (2003).
   `https://webstore.ansi.org/standards/ascx9/ansix9422003r2013`

2. *Factoring polynomials with rational coefficients.* Lenstra, A.K., & Lenstra Jr., H. W., & Lovsz, L. (1982). Mathematische Annalen, vol. 261, pp. 515-534.
   `https://infoscience.epfl.ch/record/164484/files/nscan4.PDF`

3. *Lattice basis reduction: improved practical algorithms and solving subset sum problems.* Schnorr, C.P., Euchner, M. (1994). Math. Programming 66, 181–199.
   `https://link.springer.com/article/10.1007/BF01581144`

4. *Improved algorithms for integer programming and related lattice problems.* Kannan, R. (1983). STOC 1983, pp. 193–206.
   `https://dl.acm.org/doi/10.1145/800061.808749`

5. *Efficient signature generation by smart cards.* Schnorr, C.P (1991). J. Cryptology 4, pp. 161–174.
   `https://doi.org/10.1007/BF00196725`